

# Computer Vision HW2 Report


Student ID: R12521601

Name: 詹承諺


## Part 1. (10%)

- Plot confusion matrix of two settings. (i.e. Bag of sift and tiny image) (5%)

Ans:



Bag of sift  
(accuracy: 0.654)



Tiny image  
(accuracy: 0.213)

- Compare the results/accuracy of both settings and explain the result. (5%)

Ans:

**Accuracy:**

- (1) Bag of sift: 0.654
- (2) Tiny image: 0.213

**Setting:**

**(1) Bag of sift:**

計算距離的部分在 `get_bags_of_sifts` 的 function 中 `cdist` 是用 default 的 metric ('euclidean')，在每個 function 中都有先對 image 做 normalization。

**(2) Tiny image:**

先將 image resize 成 16\*16，再將其 flatten 並 normalize。

**(3) nearest\_neighbor\_classify**

先將 image 做 normalize，在 `nearest_neighbor_classify` 的 function 中 `cdist` 的 metric 是用 'minkowski'

且  $p$  設為 0.5。KNN 的部分則是將  $k$  設置為 6。

## Result:

Bag of sift 相較於 Tiny image 來說準確率高了許多，Bag of sift 的 confusion matrix 對角線顏色明顯較深，而 Tiny image 的 confusion matrix 中則較無明顯的對角線，許多圖片都會分類成 HW, OC, CST。

## Part 2. (25%)

### • Report accuracy of both models on the validation set. (2%)

Ans:

(1) MyNet: 0.8440

(2) ResNet18: 0.9126

### • Print the network architecture & number of parameters of both models. What is the main difference between ResNet and other CNN architectures? (5%)

Ans:

#### (1) MyNet

Layer (type)	Output Shape	Param #
Conv2d-1	[-1, 64, 32, 32]	1,792
BatchNorm2d-2	[-1, 64, 32, 32]	128
ReLU-3	[-1, 64, 32, 32]	0
MaxPool2d-4	[-1, 64, 16, 16]	0
Conv2d-5	[-1, 128, 16, 16]	73,856
BatchNorm2d-6	[-1, 128, 16, 16]	256
ReLU-7	[-1, 128, 16, 16]	0
MaxPool2d-8	[-1, 128, 8, 8]	0
Conv2d-9	[-1, 256, 8, 8]	295,168
BatchNorm2d-10	[-1, 256, 8, 8]	512
ReLU-11	[-1, 256, 8, 8]	0
MaxPool2d-12	[-1, 256, 4, 4]	0
Conv2d-13	[-1, 512, 4, 4]	1,180,160
BatchNorm2d-14	[-1, 512, 4, 4]	1,024
ReLU-15	[-1, 512, 4, 4]	0
MaxPool2d-16	[-1, 512, 2, 2]	0
Linear-17	[-1, 1024]	2,098,176
ReLU-18	[-1, 1024]	0
Dropout-19	[-1, 1024]	0
Linear-20	[-1, 512]	524,800
ReLU-21	[-1, 512]	0
Dropout-22	[-1, 512]	0
Linear-23	[-1, 10]	5,130

---

---

**Total params: 4,181,002**

**Trainable params: 4,181,002**

**Non-trainable params: 0**

---

**Input size (MB): 0.01**

**Forward/backward pass size (MB): 3.08**

**Params size (MB): 15.95**

**Estimated Total Size (MB): 19.04**

MyNet(

(conv1): Conv2d(3, 64, kernel\_size=(3, 3), stride=(1, 1), padding=(1, 1))

(batchnorm1): BatchNorm2d(64, eps=1e-05, momentum=0.1, affine=True, track\_running\_stats=True)

(conv2): Conv2d(64, 128, kernel\_size=(3, 3), stride=(1, 1), padding=(1, 1))

(batchnorm2): BatchNorm2d(128, eps=1e-05, momentum=0.1, affine=True, track\_running\_stats=True)

(conv3): Conv2d(128, 256, kernel\_size=(3, 3), stride=(1, 1), padding=(1, 1))

(batchnorm3): BatchNorm2d(256, eps=1e-05, momentum=0.1, affine=True, track\_running\_stats=True)

(conv4): Conv2d(256, 512, kernel\_size=(3, 3), stride=(1, 1), padding=(1, 1))

(batchnorm4): BatchNorm2d(512, eps=1e-05, momentum=0.1, affine=True, track\_running\_stats=True)

(pool): MaxPool2d(kernel\_size=2, stride=2, padding=0, dilation=1, ceil\_mode=False)

(fc1): Linear(in\_features=2048, out\_features=1024, bias=True)

(fc2): Linear(in\_features=1024, out\_features=512, bias=True)

(fc3): Linear(in\_features=512, out\_features=10, bias=True)

(relu): ReLU()

(dropout): Dropout(p=0.5, inplace=False)

)

---

## (2) ResNet18

---

Layer (type)	Output Shape	Param #
Conv2d-1	[-1, 64, 32, 32]	1,792
BatchNorm2d-2	[-1, 64, 32, 32]	128
ReLU-3	[-1, 64, 32, 32]	0
Identity-4	[-1, 64, 32, 32]	0
Conv2d-5	[-1, 64, 32, 32]	36,864
BatchNorm2d-6	[-1, 64, 32, 32]	128
ReLU-7	[-1, 64, 32, 32]	0
Conv2d-8	[-1, 64, 32, 32]	36,864
BatchNorm2d-9	[-1, 64, 32, 32]	128
ReLU-10	[-1, 64, 32, 32]	0

<b>BasicBlock-11</b>	<b>[-1, 64, 32, 32]</b>	<b>0</b>
<b>Conv2d-12</b>	<b>[-1, 64, 32, 32]</b>	<b>36,864</b>
<b>BatchNorm2d-13</b>	<b>[-1, 64, 32, 32]</b>	<b>128</b>
<b>ReLU-14</b>	<b>[-1, 64, 32, 32]</b>	<b>0</b>
<b>Conv2d-15</b>	<b>[-1, 64, 32, 32]</b>	<b>36,864</b>
<b>BatchNorm2d-16</b>	<b>[-1, 64, 32, 32]</b>	<b>128</b>
<b>ReLU-17</b>	<b>[-1, 64, 32, 32]</b>	<b>0</b>
<b>BasicBlock-18</b>	<b>[-1, 64, 32, 32]</b>	<b>0</b>
<b>Conv2d-19</b>	<b>[-1, 128, 16, 16]</b>	<b>73,728</b>
<b>BatchNorm2d-20</b>	<b>[-1, 128, 16, 16]</b>	<b>256</b>
<b>ReLU-21</b>	<b>[-1, 128, 16, 16]</b>	<b>0</b>
<b>Conv2d-22</b>	<b>[-1, 128, 16, 16]</b>	<b>147,456</b>
<b>BatchNorm2d-23</b>	<b>[-1, 128, 16, 16]</b>	<b>256</b>
<b>Conv2d-24</b>	<b>[-1, 128, 16, 16]</b>	<b>8,192</b>
<b>BatchNorm2d-25</b>	<b>[-1, 128, 16, 16]</b>	<b>256</b>
<b>ReLU-26</b>	<b>[-1, 128, 16, 16]</b>	<b>0</b>
<b>BasicBlock-27</b>	<b>[-1, 128, 16, 16]</b>	<b>0</b>
<b>Conv2d-28</b>	<b>[-1, 128, 16, 16]</b>	<b>147,456</b>
<b>BatchNorm2d-29</b>	<b>[-1, 128, 16, 16]</b>	<b>256</b>
<b>ReLU-30</b>	<b>[-1, 128, 16, 16]</b>	<b>0</b>
<b>Conv2d-31</b>	<b>[-1, 128, 16, 16]</b>	<b>147,456</b>
<b>BatchNorm2d-32</b>	<b>[-1, 128, 16, 16]</b>	<b>256</b>
<b>ReLU-33</b>	<b>[-1, 128, 16, 16]</b>	<b>0</b>
<b>BasicBlock-34</b>	<b>[-1, 128, 16, 16]</b>	<b>0</b>
<b>Conv2d-35</b>	<b>[-1, 256, 8, 8]</b>	<b>294,912</b>
<b>BatchNorm2d-36</b>	<b>[-1, 256, 8, 8]</b>	<b>512</b>
<b>ReLU-37</b>	<b>[-1, 256, 8, 8]</b>	<b>0</b>
<b>Conv2d-38</b>	<b>[-1, 256, 8, 8]</b>	<b>589,824</b>
<b>BatchNorm2d-39</b>	<b>[-1, 256, 8, 8]</b>	<b>512</b>
<b>Conv2d-40</b>	<b>[-1, 256, 8, 8]</b>	<b>32,768</b>
<b>BatchNorm2d-41</b>	<b>[-1, 256, 8, 8]</b>	<b>512</b>
<b>ReLU-42</b>	<b>[-1, 256, 8, 8]</b>	<b>0</b>
<b>BasicBlock-43</b>	<b>[-1, 256, 8, 8]</b>	<b>0</b>
<b>Conv2d-44</b>	<b>[-1, 256, 8, 8]</b>	<b>589,824</b>
<b>BatchNorm2d-45</b>	<b>[-1, 256, 8, 8]</b>	<b>512</b>
<b>ReLU-46</b>	<b>[-1, 256, 8, 8]</b>	<b>0</b>
<b>Conv2d-47</b>	<b>[-1, 256, 8, 8]</b>	<b>589,824</b>
<b>BatchNorm2d-48</b>	<b>[-1, 256, 8, 8]</b>	<b>512</b>
<b>ReLU-49</b>	<b>[-1, 256, 8, 8]</b>	<b>0</b>
<b>BasicBlock-50</b>	<b>[-1, 256, 8, 8]</b>	<b>0</b>
<b>Conv2d-51</b>	<b>[-1, 512, 4, 4]</b>	<b>1,179,648</b>
<b>BatchNorm2d-52</b>	<b>[-1, 512, 4, 4]</b>	<b>1,024</b>

ReLU-53	[-1, 512, 4, 4]	0
Conv2d-54	[-1, 512, 4, 4]	2,359,296
BatchNorm2d-55	[-1, 512, 4, 4]	1,024
Conv2d-56	[-1, 512, 4, 4]	131,072
BatchNorm2d-57	[-1, 512, 4, 4]	1,024
ReLU-58	[-1, 512, 4, 4]	0
BasicBlock-59	[-1, 512, 4, 4]	0
Conv2d-60	[-1, 512, 4, 4]	2,359,296
BatchNorm2d-61	[-1, 512, 4, 4]	1,024
ReLU-62	[-1, 512, 4, 4]	0
Conv2d-63	[-1, 512, 4, 4]	2,359,296
BatchNorm2d-64	[-1, 512, 4, 4]	1,024
ReLU-65	[-1, 512, 4, 4]	0
BasicBlock-66	[-1, 512, 4, 4]	0
AdaptiveAvgPool2d-67	[-1, 512, 1, 1]	0
Linear-68	[-1, 10]	5,130
ResNet-69	[-1, 10]	0

---

**Total params: 11,174,026**

**Trainable params: 11,174,026**

**Non-trainable params: 0**

---

**Input size (MB): 0.01**

**Forward/backward pass size (MB): 16.00**

**Params size (MB): 42.63**

**Estimated Total Size (MB): 58.64**

---

ResNet18(

(resnet): ResNet(

(conv1): Conv2d(3, 64, kernel\_size=(3, 3), stride=(1, 1), padding=(1, 1))

(bn1): BatchNorm2d(64, eps=1e-05, momentum=0.1, affine=True, track\_running\_stats=True)

(relu): ReLU()

(maxpool): Identity()

(layer1): Sequential(

(0): BasicBlock(

(conv1): Conv2d(64, 64, kernel\_size=(3, 3), stride=(1, 1), padding=(1, 1), bias=False)

(bn1): BatchNorm2d(64, eps=1e-05, momentum=0.1, affine=True, track\_running\_stats=True)

(relu): ReLU(inplace=True)

(conv2): Conv2d(64, 64, kernel\_size=(3, 3), stride=(1, 1), padding=(1, 1), bias=False)

(bn2): BatchNorm2d(64, eps=1e-05, momentum=0.1, affine=True, track\_running\_stats=True)

)

(1): BasicBlock(

```

(conv1): Conv2d(64, 64, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1), bias=False)
(bn1): BatchNorm2d(64, eps=1e-05, momentum=0.1, affine=True, track_running_stats=True)
(rel): ReLU(inplace=True)
(conv2): Conv2d(64, 64, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1), bias=False)
(bn2): BatchNorm2d(64, eps=1e-05, momentum=0.1, affine=True, track_running_stats=True)
)
)
(layer2): Sequential(
  (0): BasicBlock(
    (conv1): Conv2d(64, 128, kernel_size=(3, 3), stride=(2, 2), padding=(1, 1), bias=False)
    (bn1): BatchNorm2d(128, eps=1e-05, momentum=0.1, affine=True, track_running_stats=True)
    (relu): ReLU(inplace=True)
    (conv2): Conv2d(128, 128, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1), bias=False)
    (bn2): BatchNorm2d(128, eps=1e-05, momentum=0.1, affine=True, track_running_stats=True)
    (downsample): Sequential(
      (0): Conv2d(64, 128, kernel_size=(1, 1), stride=(2, 2), bias=False)
      (1): BatchNorm2d(128, eps=1e-05, momentum=0.1, affine=True, track_running_stats=True)
    )
  )
  (1): BasicBlock(
    (conv1): Conv2d(128, 128, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1), bias=False)
    (bn1): BatchNorm2d(128, eps=1e-05, momentum=0.1, affine=True, track_running_stats=True)
    (relu): ReLU(inplace=True)
    (conv2): Conv2d(128, 128, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1), bias=False)
    (bn2): BatchNorm2d(128, eps=1e-05, momentum=0.1, affine=True, track_running_stats=True)
  )
)
(layer3): Sequential(
  (0): BasicBlock(
    (conv1): Conv2d(128, 256, kernel_size=(3, 3), stride=(2, 2), padding=(1, 1), bias=False)
    (bn1): BatchNorm2d(256, eps=1e-05, momentum=0.1, affine=True, track_running_stats=True)
    (relu): ReLU(inplace=True)
    (conv2): Conv2d(256, 256, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1), bias=False)
    (bn2): BatchNorm2d(256, eps=1e-05, momentum=0.1, affine=True, track_running_stats=True)
    (downsample): Sequential(
      (0): Conv2d(128, 256, kernel_size=(1, 1), stride=(2, 2), bias=False)
      (1): BatchNorm2d(256, eps=1e-05, momentum=0.1, affine=True, track_running_stats=True)
    )
  )
  (1): BasicBlock(
    (conv1): Conv2d(256, 256, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1), bias=False)
    (bn1): BatchNorm2d(256, eps=1e-05, momentum=0.1, affine=True, track_running_stats=True)

```

```

        (relu): ReLU(inplace=True)
        (conv2): Conv2d(256, 256, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1), bias=False)
        (bn2): BatchNorm2d(256, eps=1e-05, momentum=0.1, affine=True, track_running_stats=True)
    )
)
(layer4): Sequential(
  (0): BasicBlock(
    (conv1): Conv2d(256, 512, kernel_size=(3, 3), stride=(2, 2), padding=(1, 1), bias=False)
    (bn1): BatchNorm2d(512, eps=1e-05, momentum=0.1, affine=True, track_running_stats=True)
    (relu): ReLU(inplace=True)
    (conv2): Conv2d(512, 512, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1), bias=False)
    (bn2): BatchNorm2d(512, eps=1e-05, momentum=0.1, affine=True, track_running_stats=True)
    (downsample): Sequential(
      (0): Conv2d(256, 512, kernel_size=(1, 1), stride=(2, 2), bias=False)
      (1): BatchNorm2d(512, eps=1e-05, momentum=0.1, affine=True, track_running_stats=True)
    )
  )
)
  (1): BasicBlock(
    (conv1): Conv2d(512, 512, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1), bias=False)
    (bn1): BatchNorm2d(512, eps=1e-05, momentum=0.1, affine=True, track_running_stats=True)
    (relu): ReLU(inplace=True)
    (conv2): Conv2d(512, 512, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1), bias=False)
    (bn2): BatchNorm2d(512, eps=1e-05, momentum=0.1, affine=True, track_running_stats=True)
  )
)
  (avgpool): AdaptiveAvgPool2d(output_size=(1, 1))
  (fc): Linear(in_features=512, out_features=10, bias=True)
)
)

```

## Main Difference:


ResNet 中 Convolution 的 stride 有(1,1)、(2,,2)兩種，但 MyNet 只有(1,1)一種。以 Convolution 層數來說，MyNet 只有 4 層，但 ResNet18 有 18 層。ResNet18 的 Trainable Parameter 數比 MyNet 多，整體 Total size 也較大。

• Plot four learning curves (loss & accuracy) of the training process (train/validation) for both models. Total 8 plots. (8%)


Ans:

## MyNet


(1) Training Loss




(2) Training Accuracy



(3) Validation Loss




(4) Validation Accuracy




## ResNet18

(1) Training Loss




(2) Training Accuracy






### (3) Validation Loss



### (4) Validation Accuracy



• Briefly describe what method do you apply on your best model? (e.g. data augmentation, model architecture, loss function, etc) (10%)

Ans:

**Data augmentation:**

```
if split == 'train':
    transform = transforms.Compose([
        transforms.Resize((32,32)),
        ##### TODO: Data Augmentation Begin #####
        transforms.RandomHorizontalFlip(),
        transforms.RandomRotation(15),
        transforms.RandomAffine(degrees=0, translate=(0.1, 0.1)),
        transforms.ColorJitter(brightness=0.1, contrast=0.1, saturation=0.1, hue=0.1),
        transforms.RandomCrop(32, padding=3),

        ##### TODO: Data Augmentation End #####
        transforms.ToTensor(),
        transforms.Normalize(mean=[0.485, 0.456, 0.406], std=[0.229, 0.224, 0.225])
    ])
```

除了原有的 Resize 跟 Normalize 外，還有做 RandomHorizontalFlip, RandomRotation, RandomAffine, ColorJitter, RandomCrop

**Model Architecture:**

MyNet 的部分，原來前面只有四個 Convolution 並做 Maxpooling，後來將 Convolution 後面做 batch normalization 及 ReLU 準確率就有顯著的進步。

```
x = self.pool(self.relu(self.batchnorm1(self.conv1(x))))
x = self.pool(self.relu(self.batchnorm2(self.conv2(x))))
x = self.pool(self.relu(self.batchnorm3(self.conv3(x))))
x = self.pool(self.relu(self.batchnorm4(self.conv4(x))))
x = x.view(-1, 512 * 2 * 2)
x = self.dropout(self.relu(self.fc1(x)))
x = self.dropout(self.relu(self.fc2(x)))
x = self.fc3(x)
```

ResNet18 的部分，則是在後面多加一層 Convolution Layer(縮小 kernel size 至 3)及做 ReLU，並將 Maxpool Layer 換成 Identity()，最後再加一層 Fully Connected Layer。

```
# (batch_size, 3, 32, 32)
self.resnet = models.resnet18(pretrained=True)
self.resnet.conv1 = nn.Conv2d(3, 64, kernel_size=3, stride=1, padding=1)
self.resnet.relu = nn.ReLU()
self.resnet.maxpool = Identity()

self.resnet.fc = nn.Linear(self.resnet.fc.in_features, 10)
```

### Hyper Parameter (Config.py):

在 Config.py 中，Epoch 設為 100，batch size 調整成 128。Learning rate 則設為 0.01，並在 epoch 為 50, 80 的時候動態減少 learning rate。